

Aula 03

Gradiente Descendente e Otimização de Funções

Agenda

01

Recapitulação:
Função de Custo

02

Minimização da
Função de Custo

03

Gradiente
Descendente

04

Equação Normal

05

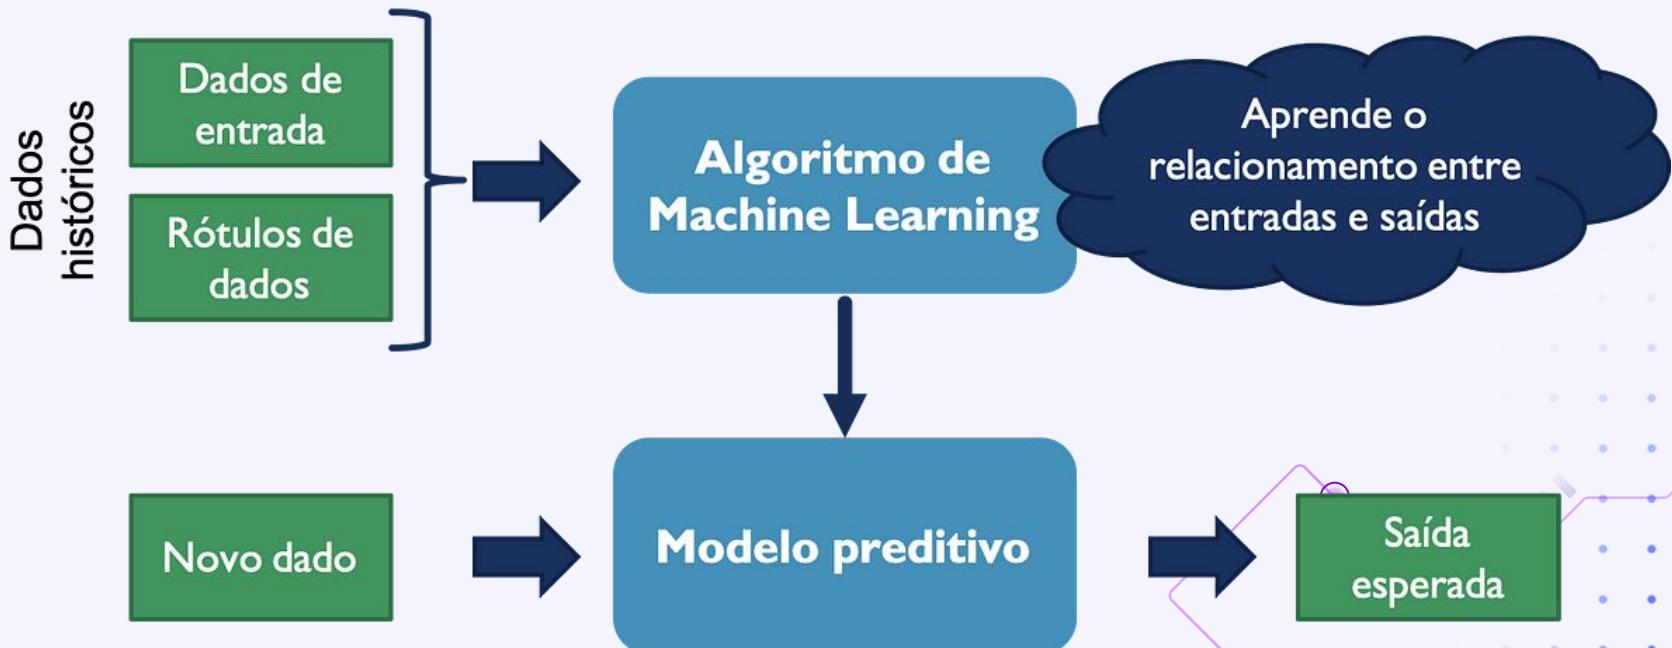
Implementando
os algoritmos

01

Recapitulação: Função de Custo



Modelando um fenômeno



Quão bom é o nosso modelo?

Como esse ajuste é calculado?

Podemos metrificar o quão bom o nosso modelo é?

Qual é o **erro** dele?

O erro do modelo é uma espécie de diferença entre o chute do modelo e a resposta correta.

Quão bom é o nosso modelo?

Como esse ajuste é calculado?

Vamos pensar no nosso erro como **um custo** ou **uma perda**.

Queremos minimizar o custo!

Objetivo: minimizar uma função que representa o custo do nosso modelo.

O que é uma Regressão Linear?

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

Uma função F é linear se:

$$F(x + y) = F(x) + F(y) \quad \text{e} \quad c * F(x) = F(c * x)$$

Pode ser representada por uma reta, um plano, ...

Qual a Função de Custo?

A nossa métrica de erro será a *Mean Squared Error*:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

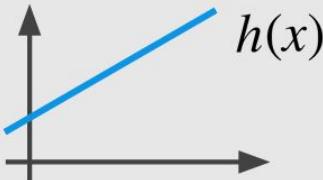
Função de Custo

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Parameters:

$$\theta_0, \theta_1$$



Cost Function:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal:

$$\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$$

Simplified

$$h_{\theta}(x) = \theta_1 x$$



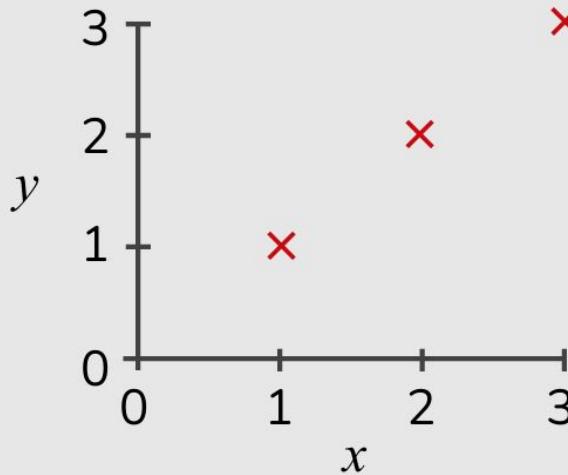
$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\underset{\theta_1}{\text{minimize}} J(\theta_1)$$

Função de Custo

$$h_{\theta}(x)$$

(for fixed θ_1 , this is a function of x)



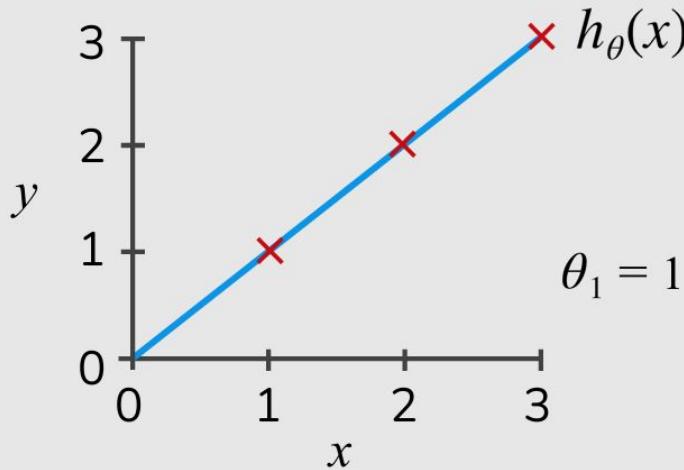
$$J(\theta_1)$$

(function of the parameters θ_1)

Função de Custo

$$h_{\theta}(x)$$

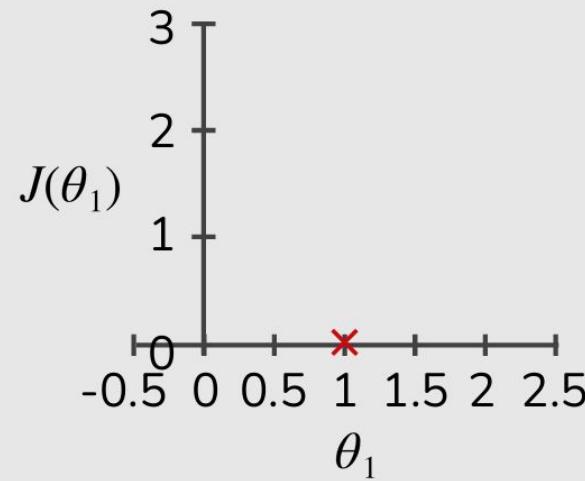
(for fixed θ_1 , this is a function of x)



$$J(\theta_1) = J(1) = 0$$

$$J(\theta_1)$$

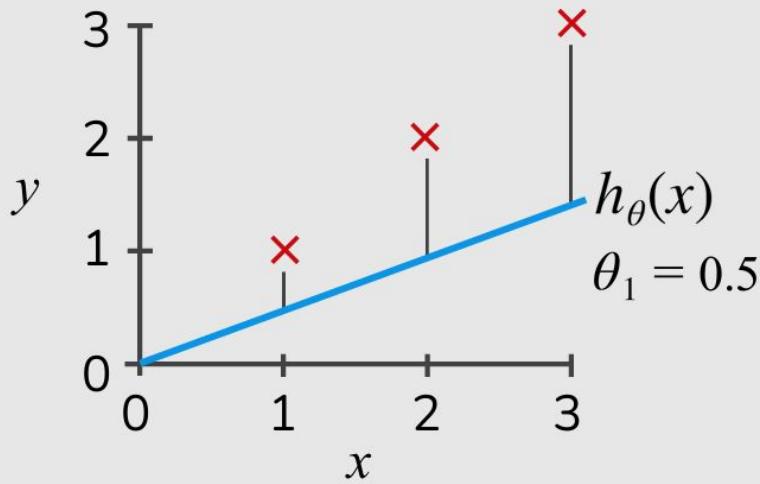
(function of the parameters θ_1)



Função de Custo

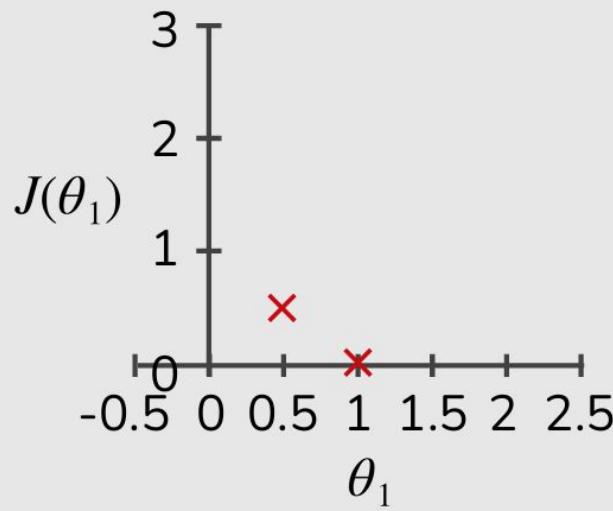
$$h_{\theta}(x)$$

(for fixed θ_1 , this is a function of x)



$$J(\theta_1)$$

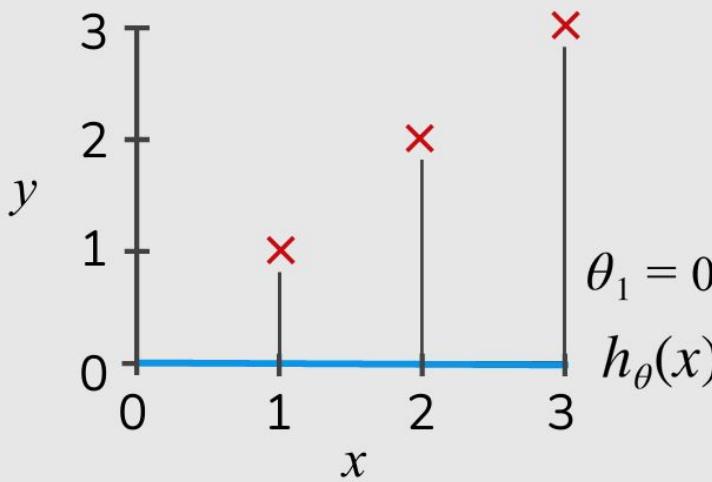
(function of the parameters θ_1)



Função de Custo

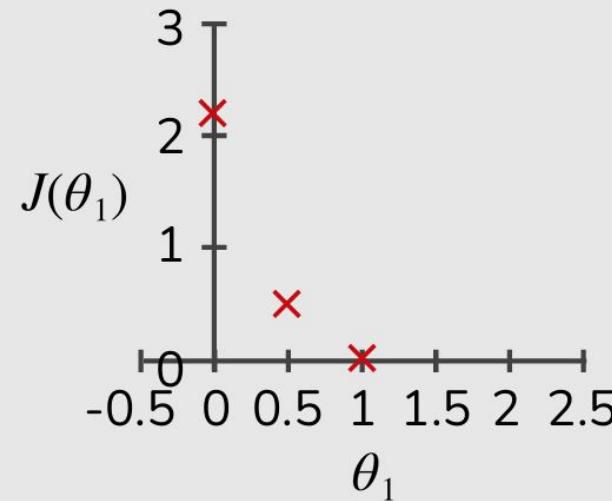
$$h_{\theta}(x)$$

(for fixed θ_1 , this is a function of x)



$$J(\theta_1)$$

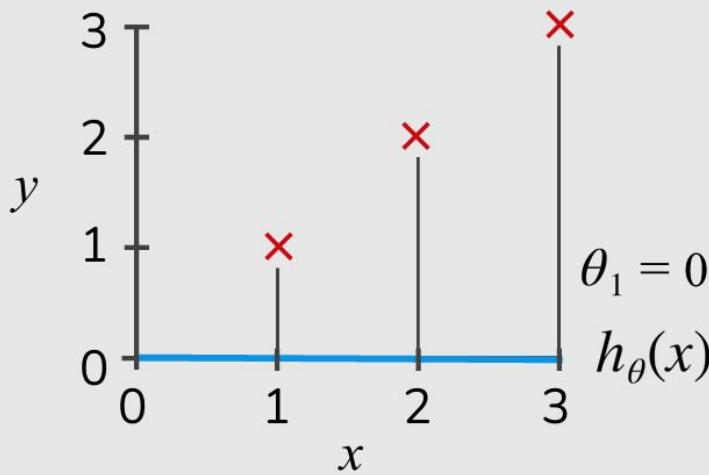
(function of the parameters θ_1)



Função de Custo

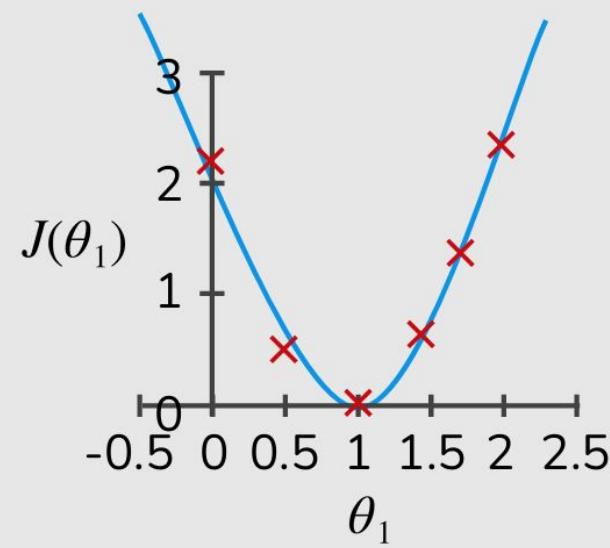
$$h_\theta(x)$$

(for fixed θ_1 , this is a function of x)



$$J(\theta_1)$$

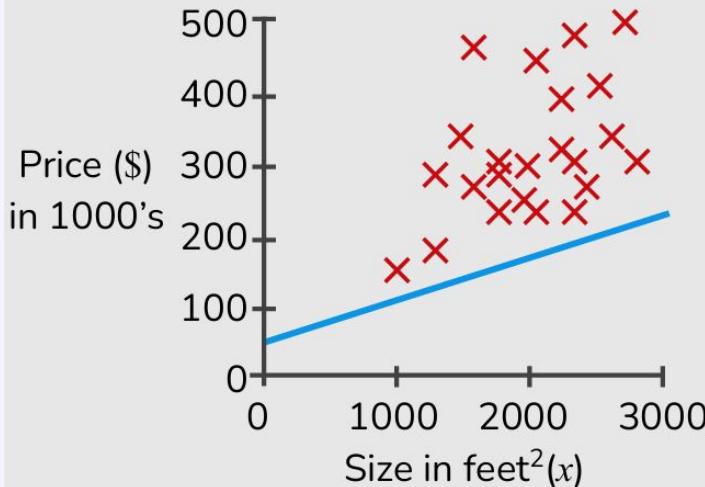
(function of the parameters θ_1)



Função de Custo

$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)

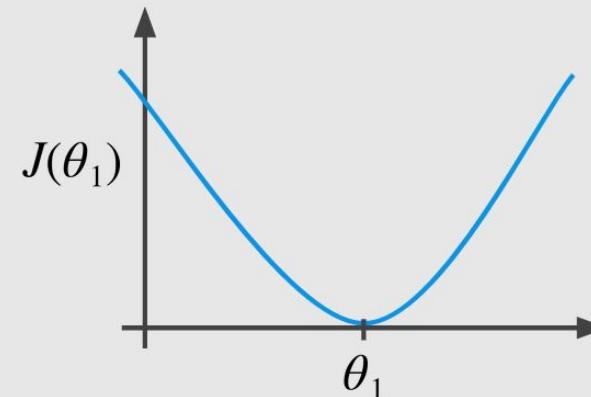


$$h_{\theta}(x) = 50 + 0.06x$$

$$\begin{aligned}\theta_0 &= 50 \\ \theta_1 &= 0.06\end{aligned}$$

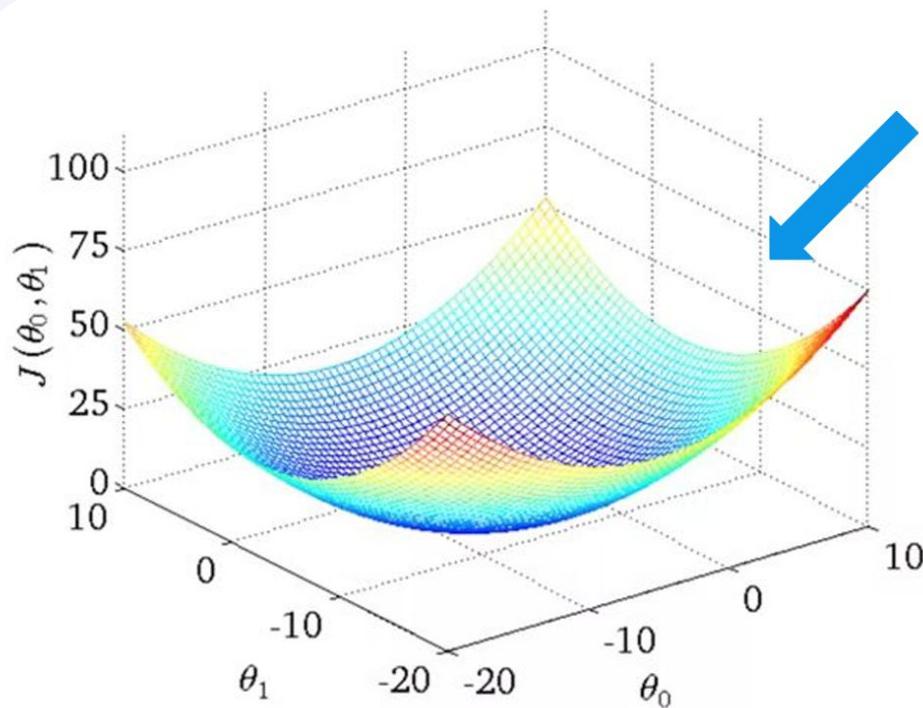
$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)



θ_0 and θ_1 ?

Função de Custo

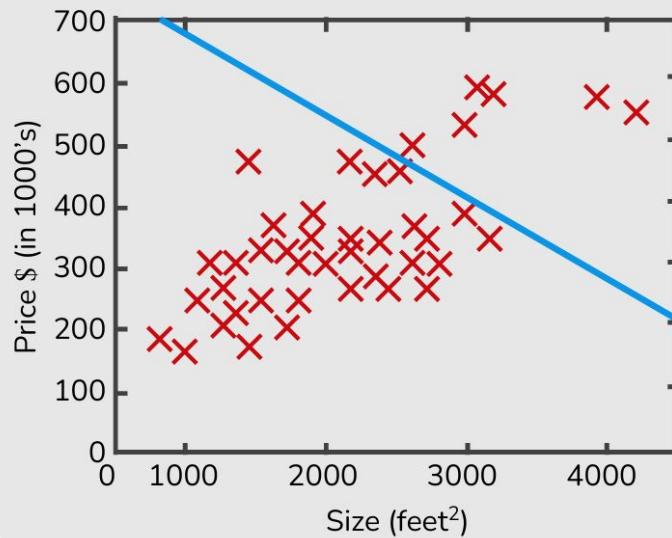


Convex
Function

Função de Custo

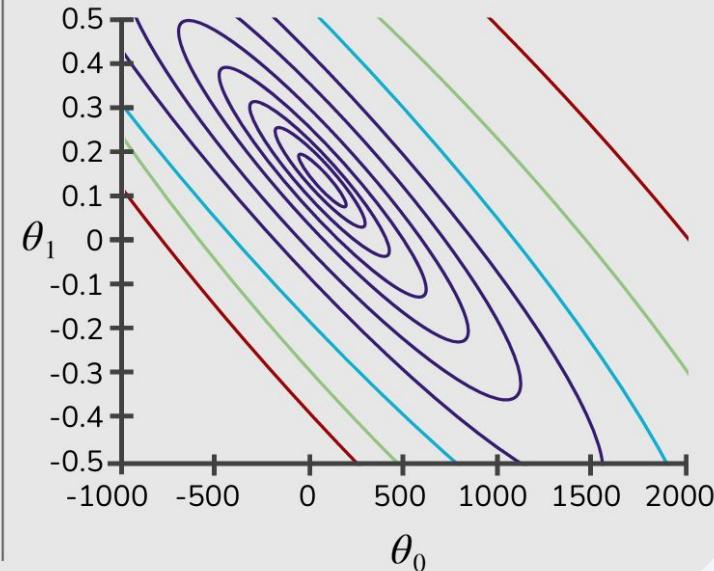
$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

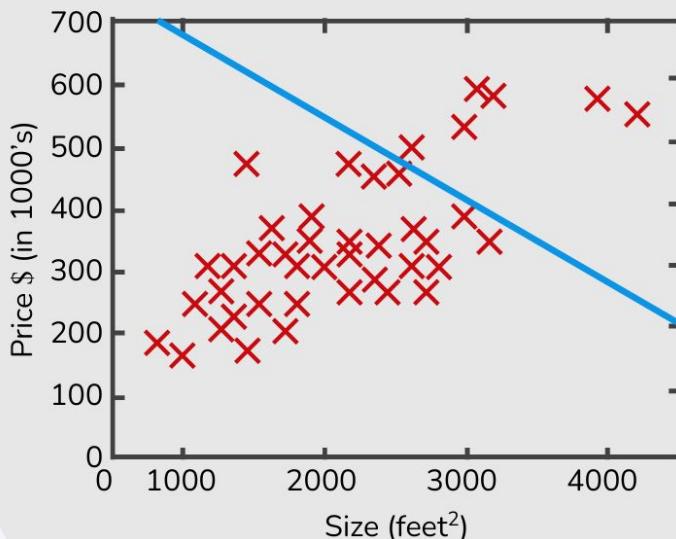
(function of the parameters θ_0, θ_1)



Função de Custo

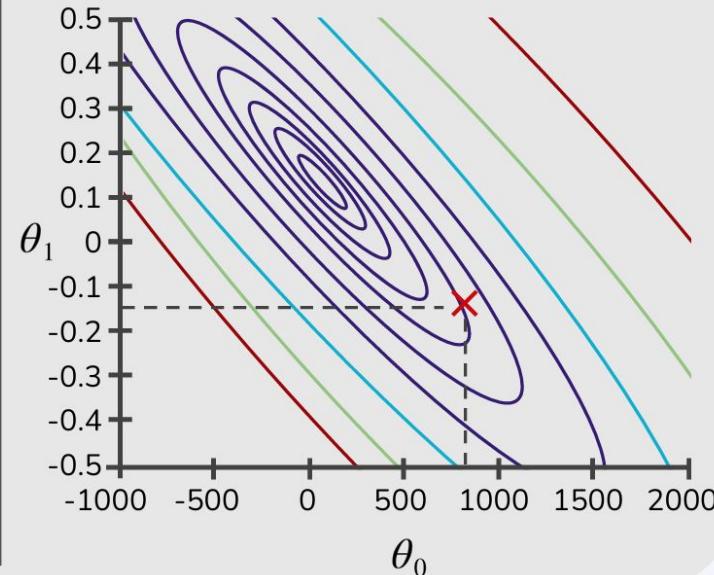
$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

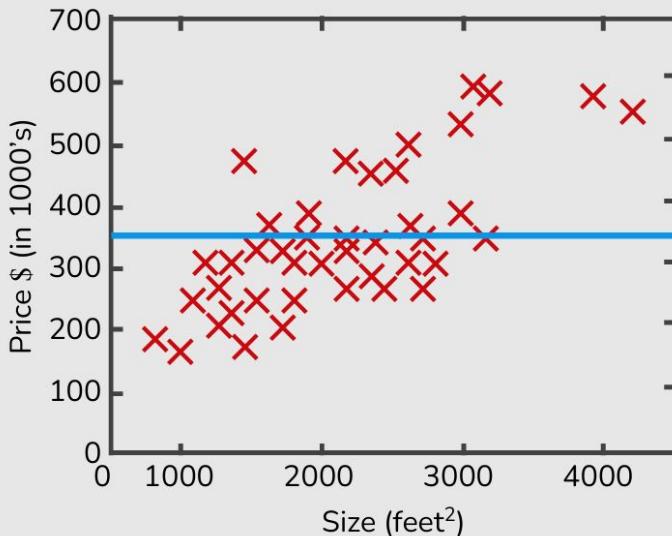
(function of the parameters θ_0, θ_1)



Função de Custo

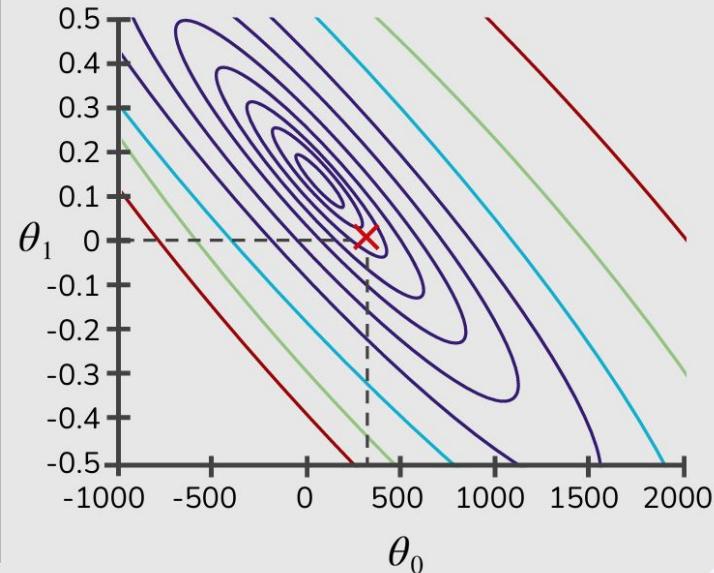
$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

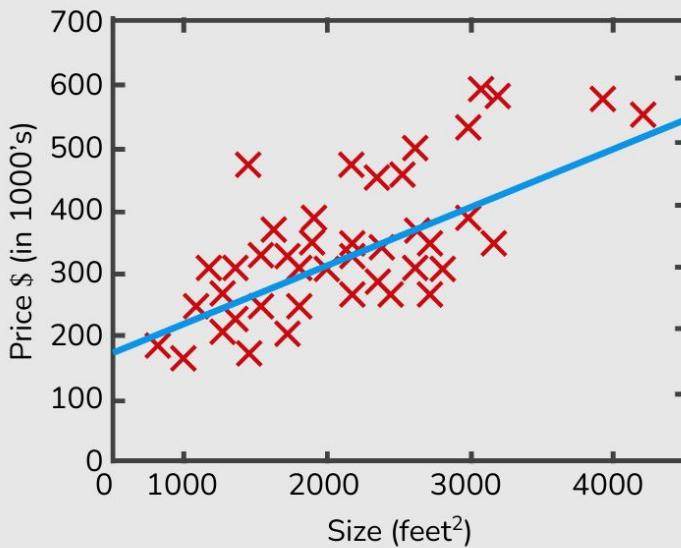
(function of the parameters θ_0, θ_1)



Função de Custo

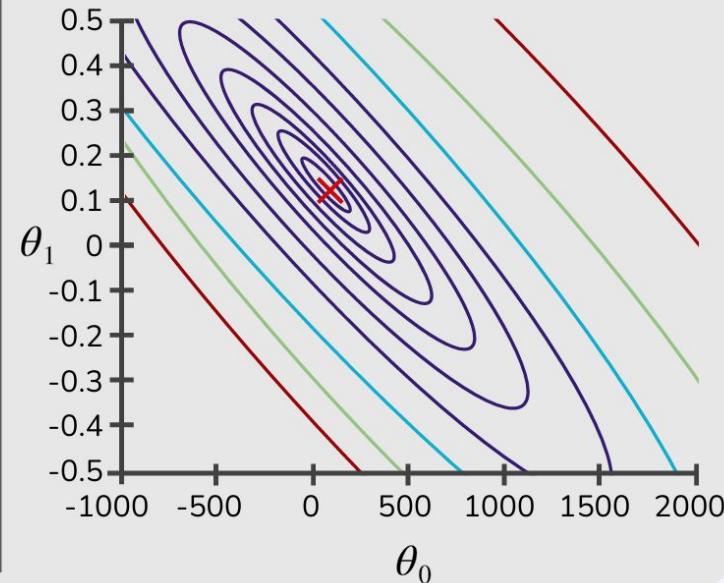
$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)



02

Minimização da Função de Custo



O que é minimizar uma função?

Minimizar uma função é encontrar os valores de entrada que retornam **o menor valor possível** para aquela função.

Este valor também é conhecido como o **mínimo global**.

Achar o mínimo global **nem sempre é possível**.

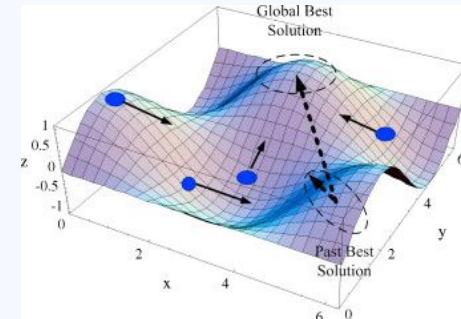
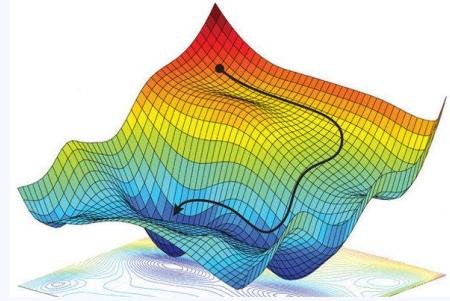
Às vezes, é possível encontrar apenas **mínimos locais**.

Encontrar **o mínimo** ou **o máximo** de funções são **problemas de otimização**.

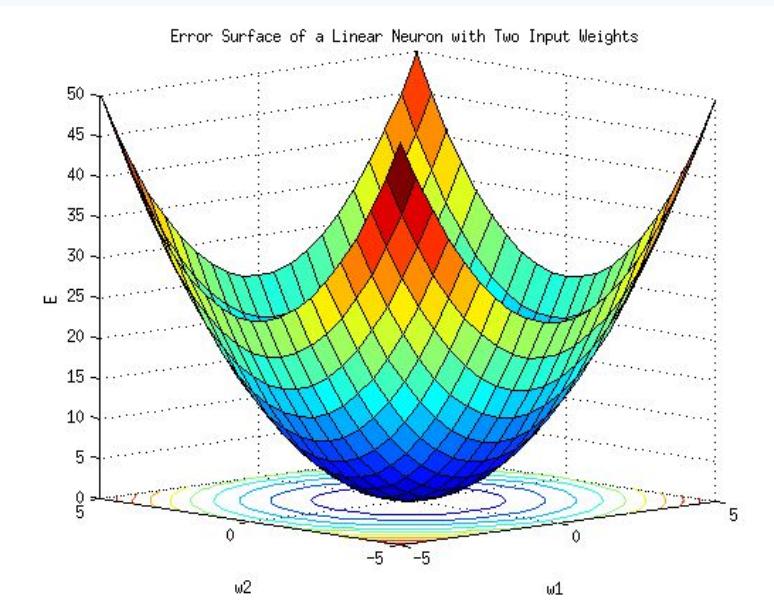
Alguns algoritmos de otimização

- MMQ → Equação Normal
(OLS) (*Normal Equation*)
- **Gradiente Descendente**
(Gradient Descent)
- Otimização por Enxame de Partículas
(*Particle Swarm Optimization*)
- Algoritmo Genético
(*Genetic Algorithm*)

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

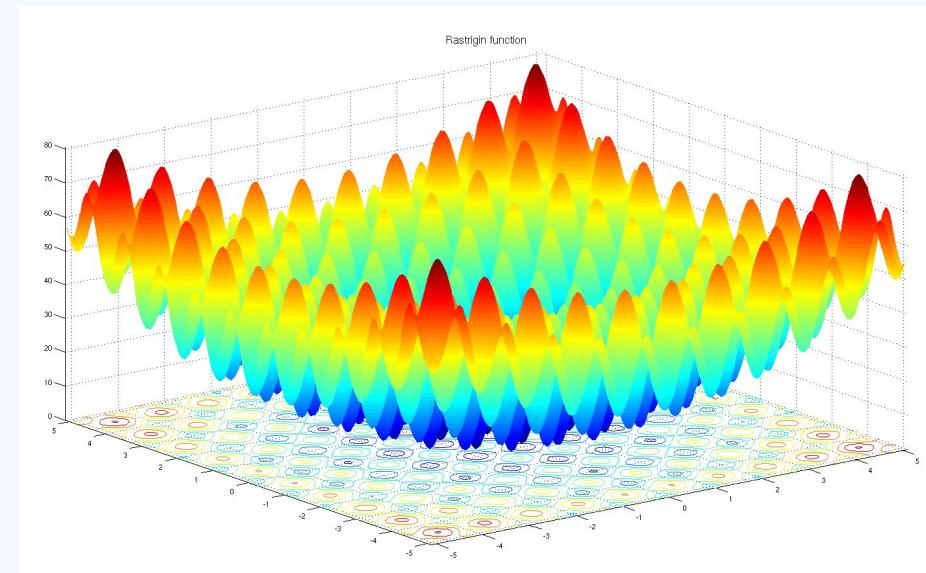


Alguns algoritmos de otimização



Função de Custo em forma de parabolóide
→ **MSE** usado em Regressão Linear

Função de benchmark (Rastrigin) de proc. otimização



Alguns algoritmos de otimização

Exemplo: a classe LinearRegression do *sklearn*.

The screenshot shows the scikit-learn API Reference for the `LinearRegression` class. The page title is "LinearRegression". The code definition is:

```
class sklearn.linear_model.LinearRegression(*, fit_intercept=True,
copy_X=True, n_jobs=None, positive=False)
```

The docstring states: "Ordinary least squares Linear Regression. LinearRegression fits a linear model with coefficients w = (w₁, ..., w_p) to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation."

Parameters:

- fit_intercept : bool, default=True**: Whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (i.e. data is expected to be centered).
- copy_X : bool, default=True**

Notes

From the implementation point of view, this is just plain Ordinary Least Squares (`scipy.linalg.lstsq`) or Non Negative Least Squares (`scipy.optimize.nnls`) wrapped as a predictor object.

Examples

```
>>> import numpy as np
>>> from sklearn.linear_model import LinearRegression
>>> X = np.array([[1, 1], [1, 2], [2, 2], [2, 3]])
>>> # y = 1 * x_0 + 2 * x_1 + 3
>>> y = np.dot(X, np.array([1, 2])) + 3
>>> reg = LinearRegression().fit(X, y)
>>> reg.score(X, y)
1.0
>>> reg.coef_
array([1., 2.])
>>> reg.intercept_
3.0...
>>> reg.predict(np.array([[3, 5]]))
array([16.])
```

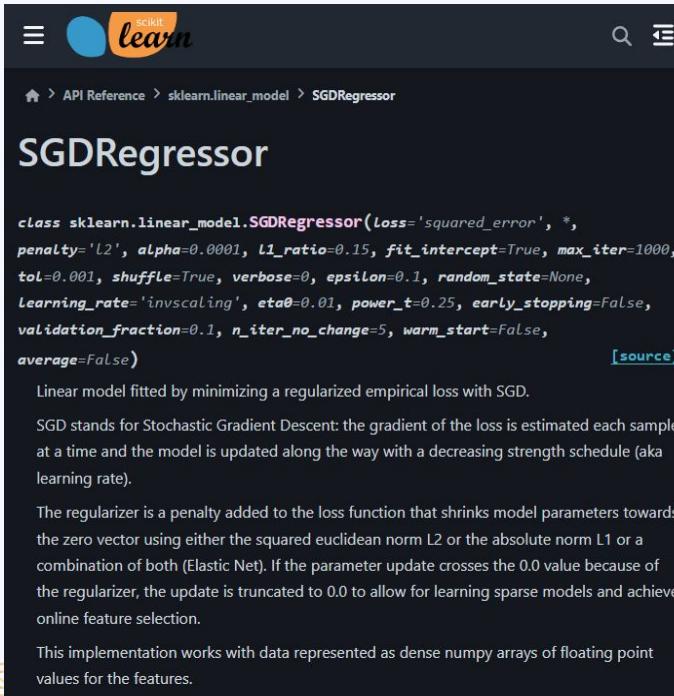
fit(x, y, sample_weight=None)

Fit linear model.

Parameters:

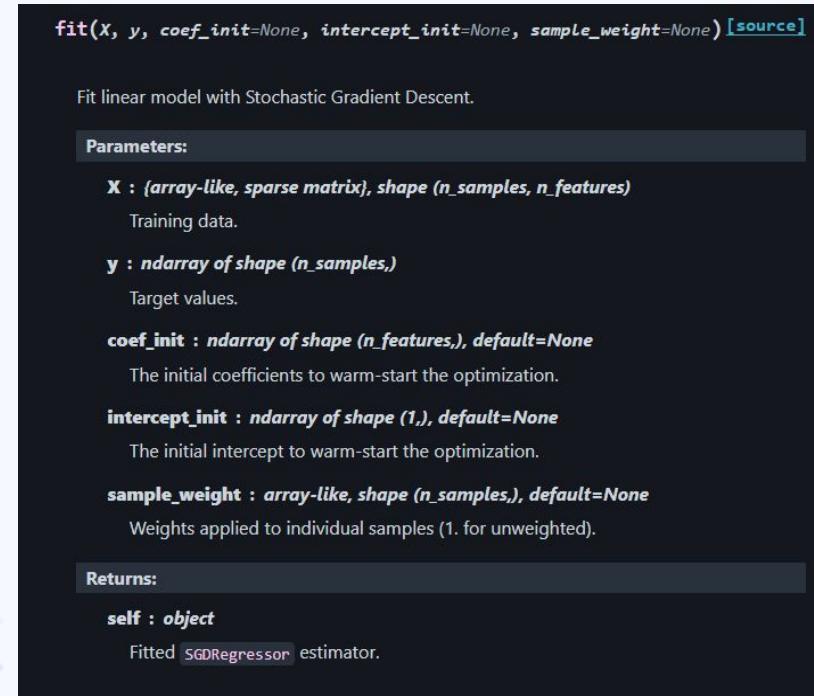
Alguns algoritmos de otimização

Exemplo: a classe SGDRegressor do *sklearn*.



The screenshot shows the scikit-learn API Reference page for the SGDRegressor class. The URL is `https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDRegressor.html`. The page title is "SGDRegressor". The code block defines the SGDRegressor class with various parameters like loss, penalty, alpha, etc. Below the code, a text block explains SGD as Stochastic Gradient Descent and its regularization mechanism. Another text block describes the regularizer as a penalty added to the loss function. At the bottom, it states that the implementation works with dense numpy arrays of floating point values for the features.

```
class sklearn.linear_model.SGDRegressor(loss='squared_error', *,\n    penalty='L2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=1000,\n    tol=0.001, shuffle=True, verbose=0, epsilon=0.1, random_state=None,\n    learning_rate='invscaling', eta0=0.01, power_t=0.25, early_stopping=False,\n    validation_fraction=0.1, n_iter_no_change=5, warm_start=False,\n    average=False)\n\n[source]\n\nLinear model fitted by minimizing a regularized empirical loss with SGD.\n\nSGD stands for Stochastic Gradient Descent: the gradient of the loss is estimated each sample at a time and the model is updated along the way with a decreasing strength schedule (aka learning rate).\n\nThe regularizer is a penalty added to the loss function that shrinks model parameters towards the zero vector using either the squared euclidean norm L2 or the absolute norm L1 or a combination of both (Elastic Net). If the parameter update crosses the 0.0 value because of the regularizer, the update is truncated to 0.0 to allow for learning sparse models and achieve online feature selection.\n\nThis implementation works with data represented as dense numpy arrays of floating point values for the features.
```



The screenshot shows the SGDRegressor API documentation from the scikit-learn source code. It includes the `fit` method signature, a description of fitting a linear model with SGD, parameter descriptions for X, y, coef_init, intercept_init, and sample_weight, and a return value description for self.

```
fit(X, y, coef_init=None, intercept_init=None, sample_weight=None) [source]\n\nFit linear model with Stochastic Gradient Descent.\n\nParameters:\n\nX : {array-like, sparse matrix}, shape (n_samples, n_features)\n    Training data.\n\ny : ndarray of shape (n_samples,)\n    Target values.\n\nc coef_init : ndarray of shape (n_features,), default=None\n    The initial coefficients to warm-start the optimization.\n\nc intercept_init : ndarray of shape (1,), default=None\n    The initial intercept to warm-start the optimization.\n\nc sample_weight : array-like, shape (n_samples,), default=None\n    Weights applied to individual samples (1. for unweighted).\n\nReturns:\n\nself : object\n    Fitted SGDRegressor estimator.
```

Alguns algoritmos de otimização

Exemplo: a classe LogisticRegression do *sklearn*.

The screenshot shows the *sklearn* API Reference page for the *LogisticRegression* class. The URL is [https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html](#). The page title is "LogisticRegression". The code block shows the class definition:

```
class sklearn.linear_model.LogisticRegression(penalty='L2', *, dual=False, tol=0.0001,
C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None,
solver='lbfgs', max_iter=100, multi_class='deprecated', verbose=0, warm_start=False,
n_jobs=None, L1_ratio=None)
```

The "source" button is highlighted. Below the code, there is a note about the classifier being a "Logistic Regression (aka logit, MaxEnt) classifier". It explains the training algorithm for multiclass cases and the solvers supported: 'lbfgs', 'liblinear', 'newton-cg', 'sag', 'saga' and 'newton-cholesky'. It also mentions that 'newton-cg', 'sag', and 'lbfgs' support L2 regularization, while 'liblinear' supports both L1 and L2 regularization.

Parameters:

- penalty: 'l1' or 'l2' (default 'l2').
- dual: boolean (default False).
- tol: float (default 0.0001).
- C: float (default 1.0).
- fit_intercept: boolean (default True).
- intercept_scaling: float (default 1).
- class_weight: dict or 'balanced' (default None).
- random_state: int, RandomState instance, or None (default None).
- solver: {'lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga'} (default 'lbfgs').
- max_iter: int (default 100).
- multi_class: {'auto', 'ovr', 'multinomial'} (default 'deprecated').
- verbose: int (default 0).
- warm_start: boolean (default False).
- n_jobs: int or None (default None).
- L1_ratio: float between 0 and 1 (default None).

The screenshot shows the detailed documentation for the *LogisticRegression* class. The "solver" parameter is highlighted. The text states that the choice of solver depends on the penalty type and the number of classes. It provides a table comparing solvers based on these factors:

solver	penalty	multinomial multiclass
'lbfgs'	'l2', None	yes
'liblinear'	'l1', 'l2'	no
'newton-cg'	'l2', None	yes
'newton-cholesky'	'l2', None	no
'sag'	'l2', None	yes
'saga'	'elasticnet', 'l1', 'l2', None	yes

Warning: The choice of the algorithm depends on the penalty chosen and on (multinomial) multiclass support.

Note: 'sag' and 'saga' fast convergence is only guaranteed on features with approximately the same scale. You can preprocess the data with a scaler from [skewer.preprocessing](#).

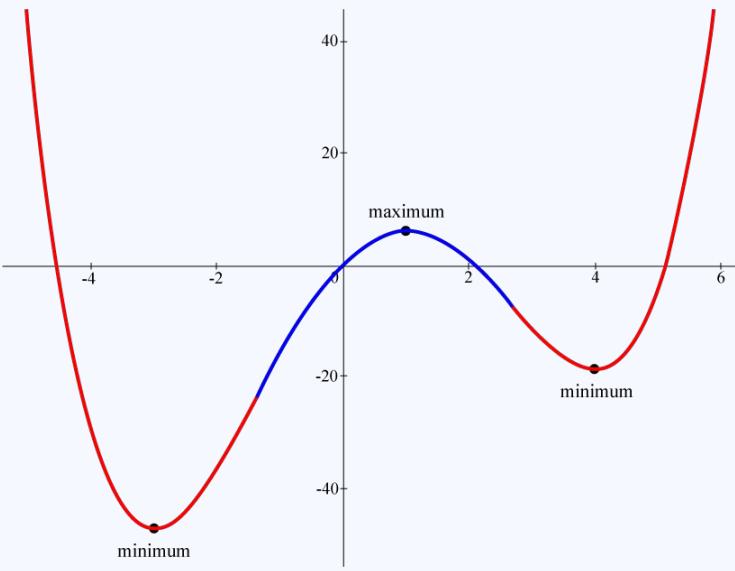
03

Gradiente Descendente

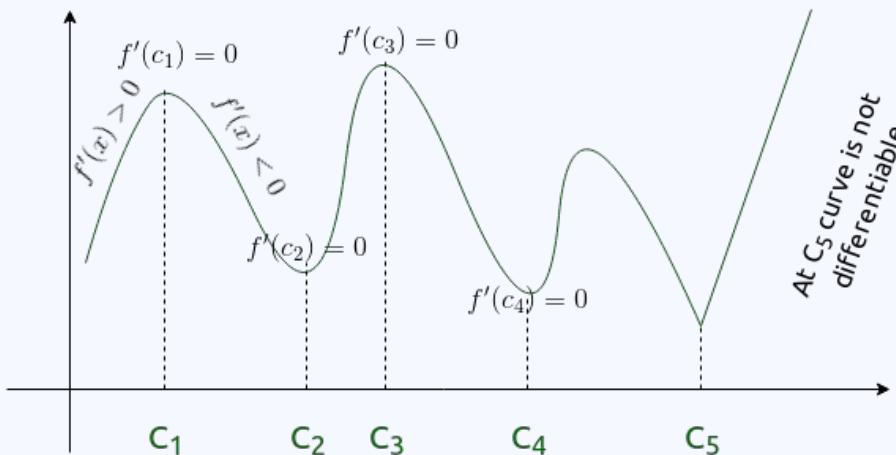


Relembrando derivadas

Como minimizamos funções em Cálculo Diferencial?



Relembrando derivadas



Derivada:

“Taxa” de variação de uma função em certo ponto.

Mínimo local:

Derivada de primeira ordem é igual a 0.

Testes de derivada de primeira e segunda ordem.

Relembrando derivadas

Vetor Gradiente:

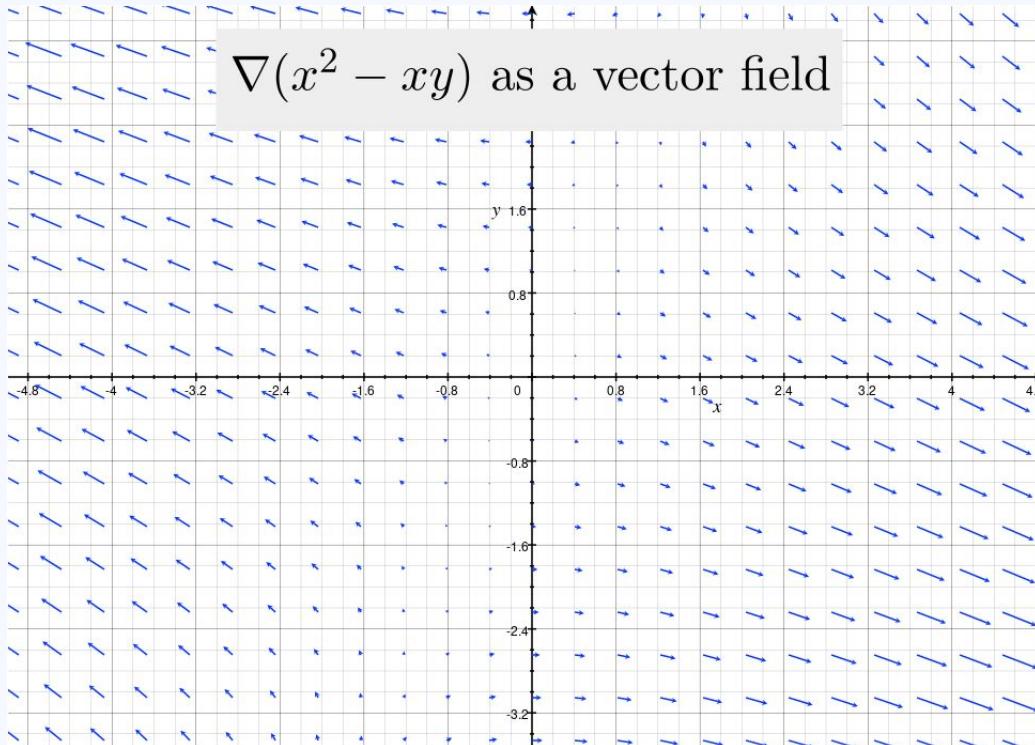
$$\nabla f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} \right)$$

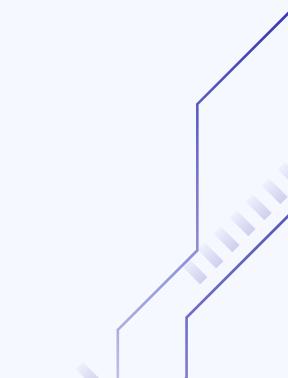
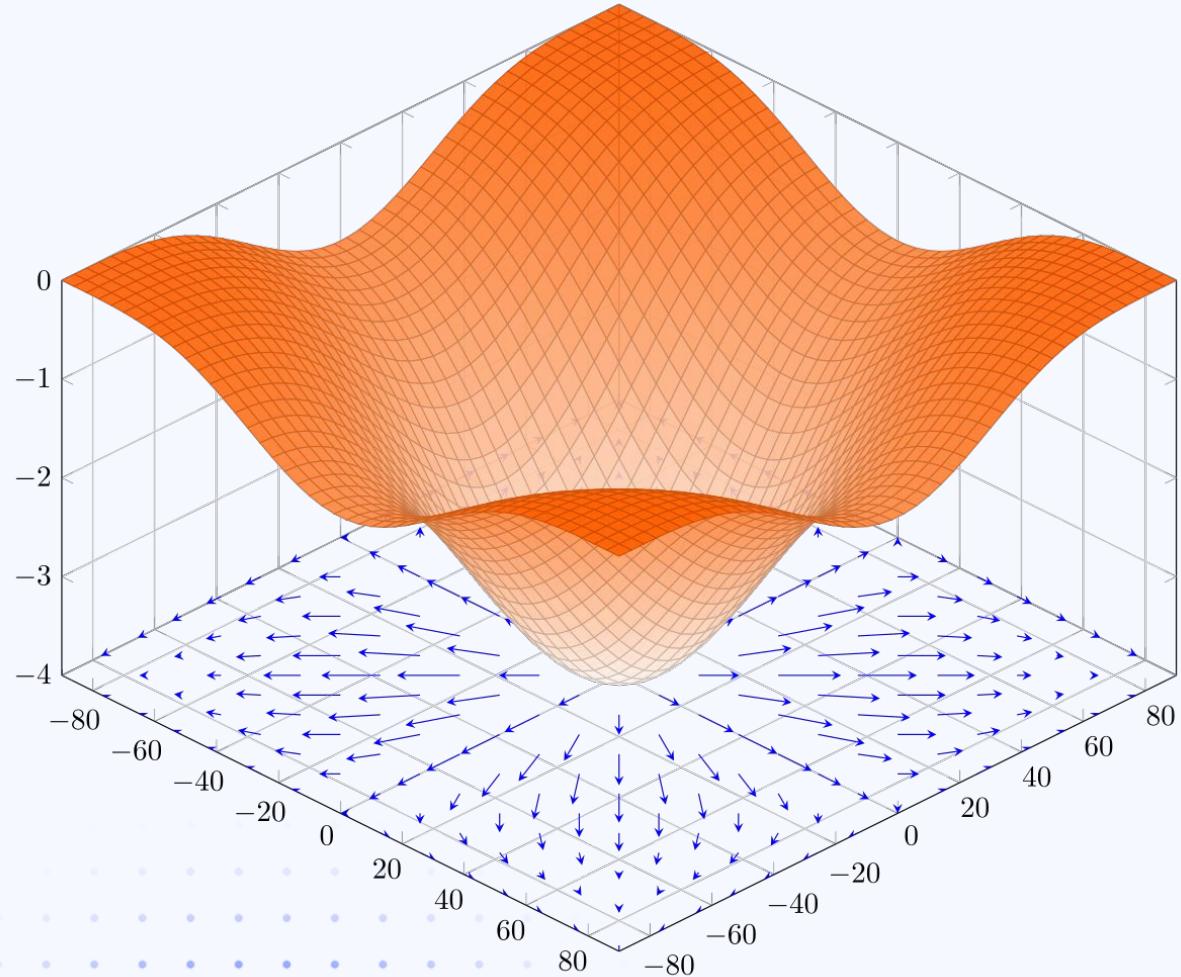
$$F(x, y, z) = x + y^2 + z^3$$

$$\nabla F(x, y, z) = \left(\frac{dF}{dx}, \frac{dF}{dy}, \frac{dF}{dz} \right) = (1, 2y, 3z^2)$$

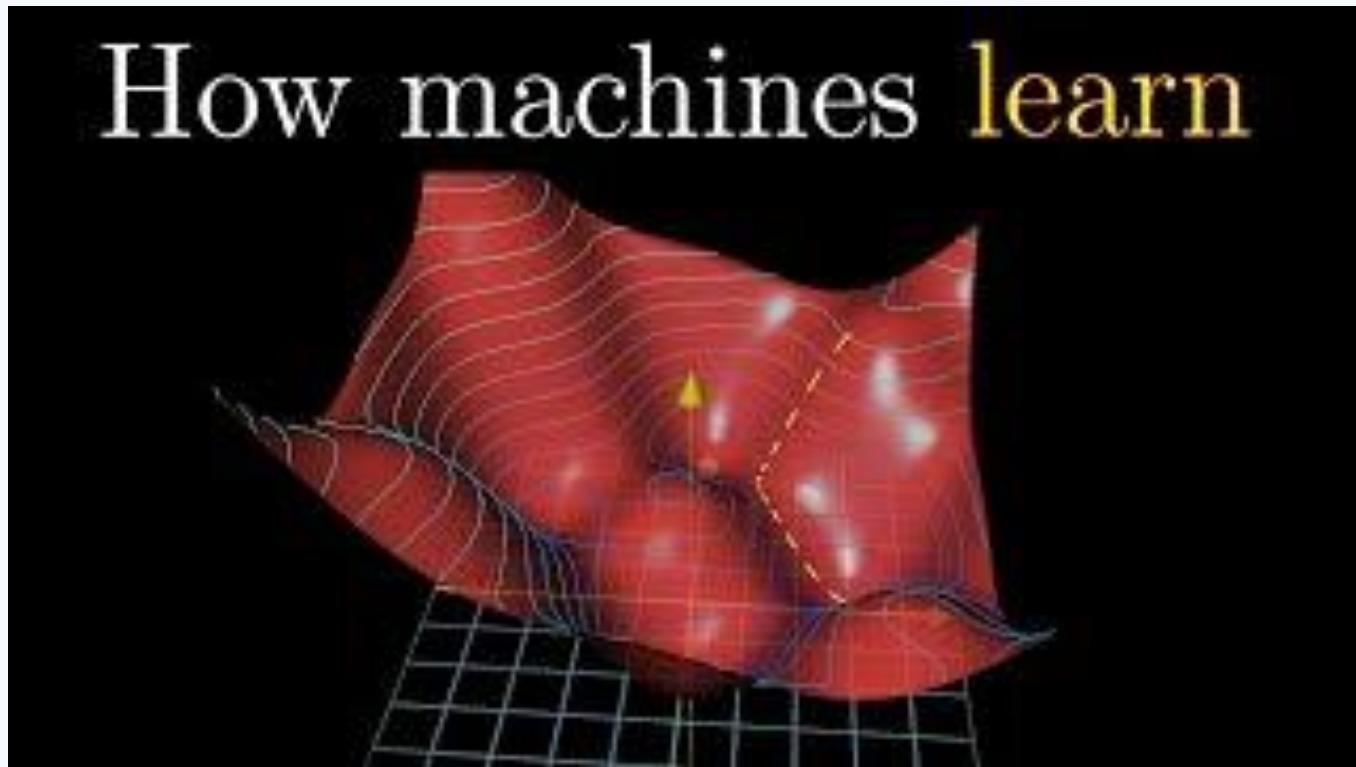
Name	Symbol	Example
Derivative	$\frac{d}{dx}$	$\frac{d}{dx}(x^2) = 2x$
Partial derivative	$\frac{\partial}{\partial x}$	$\frac{\partial}{\partial x}(x^2 - xy) = 2x - y$
Gradient	∇	$\nabla(x^2 - xy) = \begin{bmatrix} 2x - y \\ -x \end{bmatrix}$

Relembrando derivadas

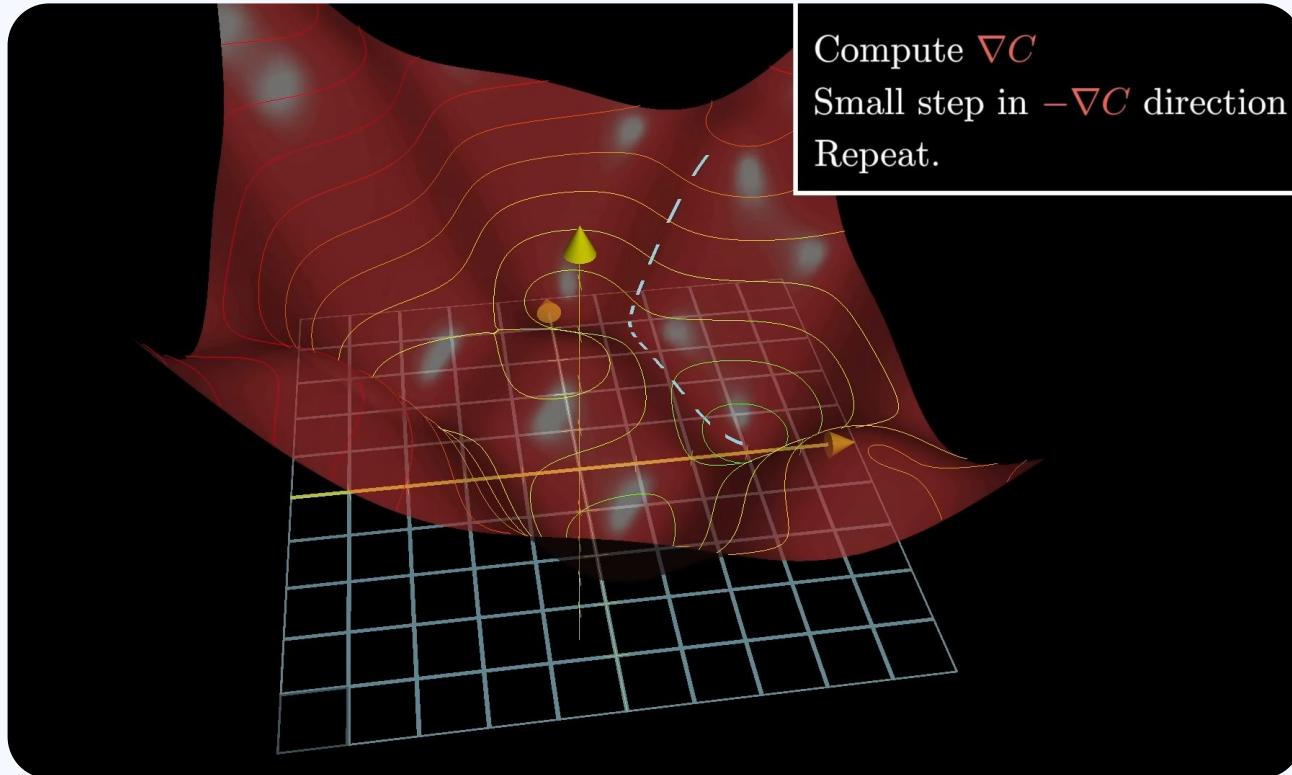




A descida do gradiente



A descida do gradiente



A descida do gradiente

$$\vec{W} = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ \vdots \\ w_{13,000} \\ w_{13,001} \\ w_{13,002} \end{bmatrix}$$

$$-\nabla C(\vec{W}) = \begin{bmatrix} 0.31 \\ 0.03 \\ -1.25 \\ \vdots \\ 0.78 \\ -0.37 \\ 0.16 \end{bmatrix}$$

w_0 should increase somewhat
 w_1 should increase a little
 w_2 should decrease a lot
 $w_{13,000}$ should increase a lot
 $w_{13,001}$ should decrease somewhat
 $w_{13,002}$ should increase a little

Treinando uma Regressão Linear

Gradient Descent algorithm

repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

(for $j = 0$ and $j = 1$)

}

Linear Regression Model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

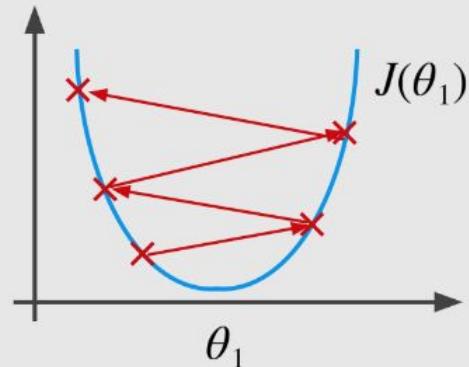
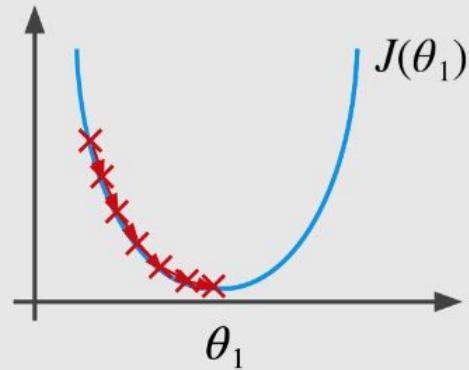
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Treinando uma Regressão Linear

$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

If α is too small, gradient descent can be slow.

If α is too large, gradient descent can be overshoot the minimum. It may fail to converge, or even diverge.



Treinando uma Regressão Linear

$$\begin{aligned}\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) &= \frac{\partial}{\partial \theta_j} \cdot \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 \\ &= \frac{\partial}{\partial \theta_j} \cdot \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2\end{aligned}$$

$$j = 0: \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})$$

$$j = 1: \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x^{(i)}$$

Treinando uma Regressão Linear

repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x^{(i)}$$

} update θ_0 and θ_1
simultaneously

}

Batch Gradient Descent

Cada **passo**, leva em conta **todas as amostras** de treinamento:

repeat until convergence {

$$\left. \begin{aligned} \theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \\ \theta_1 &:= \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x^{(i)} \end{aligned} \right\} \text{update } \theta_0 \text{ and } \theta_1 \text{ simultaneously}$$

}

Stochastic Gradient Descent

Cada **passo**, leva em conta **uma amostra** de treinamento:

```
repeat until convergence {
```

```
    for i = 1, ..., m {
```

$$\theta_0 := \theta_0 - \alpha(h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha(h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$$

```
}
```

```
}
```

Mini-batch Gradient Descent

Cada **passo**, leva em conta **b amostras** de treinamento:

Say $b = 10$, $m = 1000$.

repeat until convergence {

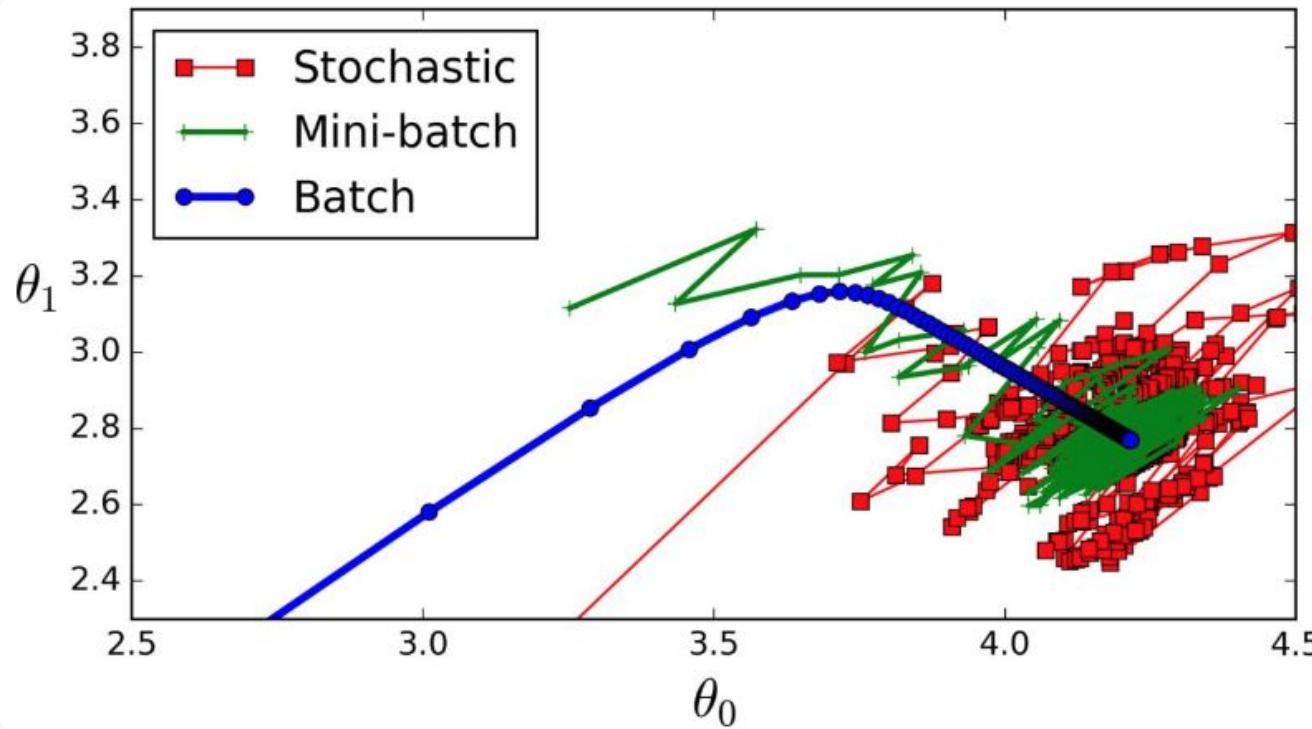
 for $i = 1, 11, 21, \dots, 991$ {

$$\theta_0 := \theta_0 - \alpha \frac{1}{10} \sum_{\substack{i=k \\ i+9}}^{i+9} (h_\theta(x^{(k)}) - y^{(k)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{10} \sum_{i=k}^{i+9} (h_\theta(x^{(k)}) - y^{(k)}) x^{(k)}$$

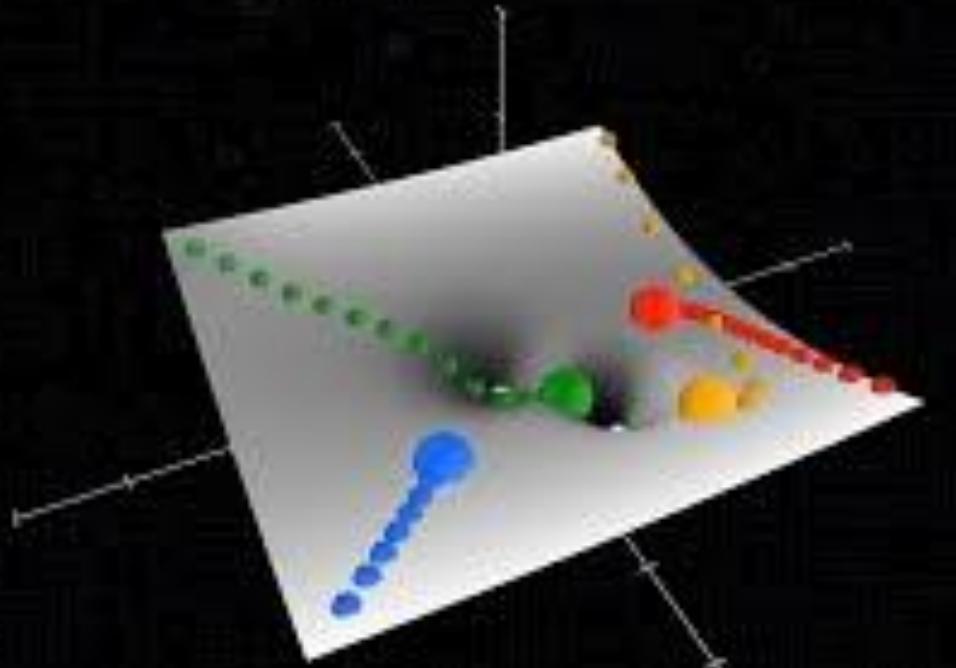
} }

Batch vs Mini-batch vs Stochastic



Optimizers in Deep Learning

- ⚙️ SGD
- ⚙️ Momentum
- ⚙️ RMSProp
- ⚙️ Adam



Instantes: 01:17-04:24; 04:38-04:54; 05:37-07:08; 12:30-12:44; 15:40-16:27; 16:28-16:50.

optimization

An overview of gradient descent optimization algorithms

Gradient descent is the preferred way to optimize neural networks and many other machine learning algorithms but is often used as a black box. This post explores how many of the most popular gradient-based optimization algorithms such as Momentum, Adagrad, and Adam actually work.

**Sebastian Ruder**

Jan 19, 2016 • 28 min read



Link para o artigo: <https://www.ruder.io/optimizing-gradient-descent/>

04

Equação Normal



O que é a Equação Normal?

Basicamente se trata de:

Implementar o Método dos Mínimos Quadrados (MMQ)

Em inglês, *Ordinary Least Squares* (*OLS*)

A **Equação Normal** representa a fórmula para resolver uma Regressão Linear **de maneira analítica**:

$$\frac{d}{d\theta} J(\theta) = \dots = 0 \quad \text{Solve for } \theta$$

Qual a Equação Normal?

Considerando **m exemplos** e **n features**, temos:

x_0	x_1	x_2	x_3	x_4	y
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}$$

Qual a Equação Normal?

Considerando **m exemplos** e **n features**, temos:

x_0	Size (feet²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$) in 1000's
	x_1	x_2	x_3	x_4	y
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix} \quad y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

Qual a Equação Normal?

Considerando **m exemplos** e **n features**, temos:

x_0	Size (feet²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$) in 1000's
	x_1	x_2	x_3	x_4	y
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix} \quad y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

Qual a Equação Normal?

Considerando **m exemplos** e **n features**, temos:

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}_{m \times (n+1)} \quad y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}_m$$

Qual a Equação Normal?

Considerando **m exemplos** e **n features**, temos:

$$x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ x_2^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$X = \begin{bmatrix} \quad (x^{(1)})^T \quad \\ \quad (x^{(2)})^T \quad \\ \vdots \\ \quad (x^{(m)})^T \quad \end{bmatrix}$$

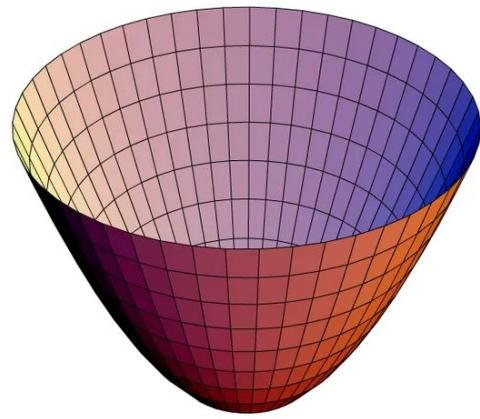
Qual a Equação Normal?

É possível chegar à seguinte expressão:

$$\theta = (X^T X)^{-1} X^T y$$

E se a matriz $X^T X$ não for invertível?

Provavelmente está utilizando duas *features* altamente correlacionadas, ou seja, **linearmente dependentes**.



https://qph.ec.quoracdn.net/main-qimg-99f7ea7d0929ed41dbecf67ec51b80b3?convert_to_webp=true

Rohan #3: Deriving the Normal Equation using matrix calculus

Understanding the analytical solution to linear regression.



Rohan Kapur · Follow

Published in A Year of Artificial Intelligence · 11 min read · Feb 16, 2016



1.2K



23



Link para o artigo:

<https://ayearofai.com/rohan-3-deriving-the-normal-equation-using-matrix-calculus-1a1b16f65dda>

05

Implementando os algoritmos





Obrigado pessoal!

Até próxima aula :)



Iris Data Science UNICAMP



@irisdatascienceunicamp

Referências

Parte do material foi adaptado dos slides da Prof^a. Sandra Avila apresentados na disciplina MC886.

Material online:

- Gradient descent, how neural networks learn | Chapter 2, Deep learning (YouTube)

Livros utilizados:

- An Introduction to Statistical Learning (2023)
- Hands-On Machine Learning with Scikit-Learn, Keras, and Tensor Flow (2017)